

# Jordan National Policy Intelligence Platform

## *Government-Grade Multi-Agent AI System – Team Guide*

### Objective

You are building a production-ready, multi-agent national intelligence platform that analyzes, monitors, and explains Jordan's strategic policies and modernization programs. This is not a prototype. This is a modular, microservice-based enterprise AI system.

### Architectural Mandate (Applies to All Teams)

- Each team owns one independent microservice.
- Each microservice exposes well-defined APIs.
- Each feature must be built as a self-contained, reusable module.
- Modules must be reusable in other applications.
- Every module must be documented.
- System must be modular, replaceable, service-oriented, independently deployable, and observable.
- Backend: FastAPI.
- Frontend: React + TypeScript + Tailwind.

### Repository & Branch Governance

Single shared GitHub repository.

Structure:

/services  
/knowledge-service  
/agent-service  
/governance-service  
/workflow-service  
/frontend-app

/shared  
/docs

Branches:

main → production-ready

dev → integration

service/\* → microservice development

feature/\* → feature branches

Only the Governance Team merges into main.

## **Team 1 – Knowledge & Retrieval Microservice (knowledge-service)**

- Mission: Build a reusable and scalable national knowledge layer.
- DNA: Data-first, structured, precise, stability-oriented.
- Modules: Ingestion engine, chunking engine, multilingual embeddings, metadata tagging, version tracking, vector retrieval API, source filtering, evidence ranking.
- APIs: /ingest, /retrieve, /sources, /versions, /health.
- Must operate independently and serve other AI systems.

## **Team 2 – Agent Orchestration Microservice (agent-service)**

- Mission: Design and implement the multi-agent orchestration layer.
- DNA: System designers, logical modelers, delegation-focused.
- Modules: Router agent, specialist agents, tool abstraction, confidence scoring, delegation engine, escalation triggers, self-verification, citation enforcement.
- APIs: /query, /delegate, /confidence, /validate, /explain\_decision.
- Must operate independently from frontend.

## **Team 3 – Governance, DevOps & Platform Integrity (governance-service)**

- Mission: Ensure quality, evaluation, monitoring, compliance, and release control.
- DNA: Institutional mindset, disciplined, risk-aware.
- Modules: Guardrails, policy enforcement, logging & audit, evaluation framework, performance metrics, CI/CD, security validation, release pipeline, benchmarking.
- APIs: /audit, /evaluate, /logs, /guardrail\_check, /metrics, /release\_status.
- This team governs merges into main.

## **Team 4 – Workflow, HITL & Institutional Operations (workflow-service)**

- Mission: Build the operational backbone of the platform.
- DNA: Process architects, reliability-focused, state-machine thinkers.
- HITL Modules: Ticket creation, escalation state machine, assignment workflow, resolution engine, human response storage.
- Operations Modules: Notification engine, conflict resolution, clarification module, feedback ingestion, response archive, SLA tracking, RBAC, admin APIs, queue prioritization, activity timeline.
- APIs: /tickets, /escalate, /resolve, /feedback, /notifications, /admin, /roles, /timeline.

## **Team 5 – Intelligence Experience & Interface Platform (frontend-app)**

- Tech Stack: React, TypeScript, Tailwind.
- Mission: Build a government-grade AI intelligence interface.
- DNA: UX clarity obsessed, trust-first design, structured interaction thinkers.
- Core UI Modules: Chat, source panel, citation viewer, confidence explanation, escalation status.
- Intelligence Features: Policy comparison, change dashboard, executive briefing UI, reading path explorer, structured answer views, export tools, role switching.
- Operational UI: Human review console, ticket dashboard, audit viewer, system metrics dashboard, notification center.
- All components must be reusable across applications.

## **System Characteristics**

- Fully modular.
- Service-isolated.
- Reusable feature components.
- Microservice-based.
- Audit-ready.
- Deployment-ready.

- API-first.
- Observability built-in.

## **What This Project Truly Becomes**

This is not a chatbot. This is not a student experiment. This is a modular, agentic, government-ready AI platform built with real service separation, DevOps discipline, institutional workflow, human override mechanisms, governance enforcement, and a production-grade frontend.